Lasso Regression

Lasso (Least Absolute Shrinkage and Selection Operator) is a type of regression that performs both variable selection and regularization in order to enhance the prediction accuracy and interpretability of the statistical model it produces. It does this by adding a penalty equivalent to the absolute value of the magnitude of the coefficients to the ordinary least squares (OLS) loss function. This penalty can shrink some coefficients to exactly zero, effectively performing feature selection by excluding those variables from the model.

Shrinkage and Selection

Lasso Regression, which stands for **Least Absolute Shrinkage and Selection Operator**, is a popular regularization technique used in linear regression. It's particularly valuable in machine learning for two main reasons:

1. **Regularization (Preventing Overfitting):** Like other regularization methods, Lasso aims to prevent overfitting. Overfitting occurs when a model learns the training data too well, including its noise and idiosyncrasies, leading to poor performance on new, unseen data. Lasso addresses this by adding a penalty term to the traditional linear regression cost function.

2. **Feature Selection:** This is a key distinguishing characteristic of Lasso. It can automatically perform feature selection by driving the coefficients of less important (or irrelevant) features to exactly zero. This effectively removes those features from the model, leading to a simpler, more interpretable, and often more robust model.

How it Works (Mathematical Intuition)

In standard linear regression (Ordinary Least Squares or OLS), the goal is to find the coefficients (β) that minimize the Residual Sum of Squares (RSS), which is the sum of the squared differences between the predicted and actual values:

$$RSS = \sum_{i=1}^{n}(y_i - \hat{y}_i)^2 = \sum_{i=1}^{n}(y_i - (\beta_0 + \sum_{j=1}^{p} x_{ij}\beta_j))^2$$

Lasso regression modifies this objective function by adding an **L1 penalty term**:

$$\text{Minimize} \quad RSS + \lambda \sum_{j=1}^{p} |\beta_j|$$

Let's break down this formula:

- RSS: This is the familiar Residual Sum of Squares, measuring how well the model fits the data.
- λ **(lambda)**: This is the **regularization parameter** (a non-negative hyperparameter). It controls the strength of the penalty.
  - **If** $\lambda = 0$: The penalty term disappears, and Lasso regression becomes equivalent to standard OLS linear regression.
  - **As** λ **increases**: The penalty for large coefficients becomes stronger, pushing them towards zero.
- $\sum_{j=1}^{p} |\beta_j|$: This is the **L1 norm** of the coefficient vector (excluding the intercept β_0). It's the sum of the absolute values of the regression coefficients. This is where the magic of feature selection happens.

The L1 Penalty and Feature Selection

The L1 penalty, due to its absolute value nature, has a unique effect. When minimizing the objective function, the model tries to keep the sum of the absolute coefficients small. To achieve this, it will tend to "prefer" setting some coefficients exactly to zero rather than shrinking all of them slightly.

Think of it geometrically:

- The OLS solution corresponds to minimizing the RSS without any constraints.
- The L1 penalty imposes a diamond-shaped constraint region around the origin in the coefficient space.
- The optimal solution for Lasso often occurs at the "corners" of this diamond, where one or more coefficients are precisely zero.

This is in contrast to **Ridge Regression (L2 regularization)**, which uses a squared penalty term ($\sum \beta_j^2$) and creates a circular constraint region. Ridge regression shrinks coefficients towards zero but rarely sets them exactly to zero.

Advantages of Lasso Regression:

- **Automatic Feature Selection:** This is its most significant advantage, especially for datasets with many features where only a subset might be truly relevant. It simplifies the model and makes it more interpretable.
- **Reduces Overfitting:** By penalizing large coefficients, it prevents the model from becoming too complex and fitting the noise in the training data.
- **Improved Interpretability:** With fewer non-zero coefficients, the model is easier to understand, as it focuses on the most impactful features.
- **Handles Multicollinearity (to some extent):** While not as robust as Elastic Net for highly correlated features, Lasso can help by selecting one feature from a group of highly correlated ones and setting the others to zero.

Limitations of Lasso Regression:

- **Bias:** Introducing a penalty term adds bias to the coefficient estimates. If the true underlying model has many small, non-zero coefficients, Lasso might inappropriately shrink some of them to zero, leading to underfitting.
- **Arbitrary Selection with Correlated Features:** If there are highly correlated features, Lasso tends to pick only one of them and set the others to zero, which might not always be ideal if all correlated features are genuinely important.
- **Cannot select more features than observations:** In cases where the number of features (p) is greater than the number of observations (n), Lasso can only select at most n features.

Hyperparameter Tuning (λ)

The choice of λ is crucial. A common approach to finding the optimal λ is through **cross-validation**. This involves:

1. Splitting the data into multiple folds.
2. Training the Lasso model on a subset of the folds for different λ values.
3. Evaluating the model's performance (e.g., Mean Squared Error) on the held-out fold.
4. Selecting the λ that minimizes the chosen evaluation metric.

In summary, Lasso Regression is a powerful tool in machine learning for building more robust, interpretable, and less prone-to-overfitting linear models by simultaneously performing regularization and automatic feature selection.

Bias and Variance of the Lasso

- **Bias:**
  - **Increases as λ increases.** A larger λ forces more coefficients towards zero, potentially moving the estimate further from the true parameters.
  - **At $\lambda = 0$:** Bias is **0** (Lasso becomes equivalent to unbiased OLS).
  - **At $\lambda = \infty$:** Bias is **high** (all coefficients are forced to zero, leading to a potentially very inaccurate model unless the true coefficients are indeed all zero).
- **Variance:**
  - **Decreases as λ increases.** A larger λ shrinks coefficients, making the model less sensitive to fluctuations in the training data.
  - **At $\lambda = 0$:** Variance is **high** (Lasso behaves like OLS, which has higher variance compared to regularized models).
  - **At $\lambda = \infty$:** Variance is **0** (all coefficients are forced to zero, resulting in a constant prediction, which by definition has no variance).

In [ ]: